

Introduction

In this notebook, we'll implement simple RNNs and LSTMs, then explore how gradients flow through these different networks.

This notebook does not require a Colab GPU. If it's enabled, you can turn it off through Runtime -> Change runtime type. (This will make it more likely for you to get Colab GPU access later in the REAL_RNN_LSTM.ipynb problem.)

Imports ¶

Note: the ipynpl installation will require you to restart the colab runtime.

```
In [ ]: ! pip install ipynpl
```

```
In [1]: import os
os.environ["KMP_DUPLICATE_LIB_OK"]="TRUE"
```

```
In [2]: import copy

# If you are not using colab you can delete these two lines
#from google.colab import output
#output.enable_custom_widget_manager()

import torch as th
from torch import nn
import torch.nn.functional as F
import torch.optim as optim
import numpy as np
import matplotlib.pyplot as plt
from ipywidgets import interactive, widgets, Layout
```

1.A: implementing a RNN layer

Consider using Pytorch's [nn.Linear](https://pytorch.org/docs/stable/generated/torch.nn.Linear.html#torch.nn.Linear) (<https://pytorch.org/docs/stable/generated/torch.nn.Linear.html#torch.nn.Linear>). You can implement this with either one Linear layer or two. If you use two, remember that you only need to include a bias term for one of the linear layers.


```

In [3]: class RNNLayer(nn.Module):
def __init__(self, input_size, hidden_size, nonlinearity=th.tanh):
    """
    Initialize a single RNN layer.

    Inputs:
    - input_size: Data input feature dimension
    - hidden_size: RNN hidden state size (also the output feature dimension)
    - nonlinearity: Nonlinearity applied to the rnn output
    """
    super().__init__()
    self.input_size = input_size
    self.hidden_size = hidden_size
    self.nonlinearity = nonlinearity
    #####
    # TODO: Initialize any parameters your class needs.
    #  $h_t = \sigma(W_h * h_{t-1} + W_x * X_t + \text{bias})$ 
    self.mixed_w = nn.Linear(input_size + hidden_size, hidden_size, bias = True)
    #####

    #####
    #                                     END OF YOUR CODE                                     #
    #####

def forward(self, x):
    """
    RNN forward pass

    Inputs:
    - x: input tensor (B, seq_len, input_size)

    Returns:
    - all_h: tensor of size (B, seq_len, hidden_size) containing hidden states
              produced for each timestep
    - last_h: hidden state from the last timestep (B, hidden_size)
    """
    h_list = [] # List to store the hidden states [h_1, ... h_T]
    #####
    # TODO: Implement the RNN forward step                                     #
    # 1. Initialize h0 with zeros                                           #
    # 2. Roll out the RNN over the sequence, storing hidden states in h_list #
    # 3. Return the appropriate outputs                                     #
    #####

    batch_size, seq_len = x.shape[:2]

    begin_pad = th.zeros((batch_size, self.hidden_size)).float()
    h_i = begin_pad

    for i in range(seq_len):
        x_i = x[:, i]
        inputs = th.cat([x_i, h_i], dim=1)
        h_i = self.nonlinearity(self.mixed_w(inputs))
        h_list.append(h_i)

    last_h = h_i

    #####
    #                                     END OF YOUR CODE                                     #
    #####

```

```

# h_list should now contain all hidden states, each of size (B, hidden_size)
# We will store the hidden states so we can analyze their gradients later
self.store_h_for_grad(h_list)

print("batch_size: ", batch_size)
print("seq_len: ", seq_len)
print("hidden_size: ", self.hidden_size)

print(len(h_list), h_list[0].shape)
all_h = th.stack(h_list, dim=1)
print(all_h.shape)
return all_h, last_h

def store_h_for_grad(self, h_list):
    """
    Store input list and allow gradient computation for all list elements
    """
    for h in h_list:
        h.retain_grad()
    self.h_list = h_list

```

Test Cases

If your implementation is correct, you should expect to see errors of less than $1e-4$.

```

In [4]: rnn = RNNLayer(1, 1)
# Overwrite initial parameters with fixed values.
# Should give deterministic results even with different implementations.
rnn.load_state_dict({k: v * 0 + .1 for k, v in rnn.state_dict().items()})
data = th.ones((1, 1, 1))
expected_out = th.FloatTensor([[[0.1973753273487091]]])
all_h, last_h = rnn(data)
assert all_h.shape == expected_out.shape
assert th.all(th.isclose(all_h, last_h))
print(f'Expected: {expected_out.item()}, got: {last_h.item()}, max error: {th.max(th

rnn = RNNLayer(2, 3, nonlinearity=lambda x: x) # no nonlinearity

num_params = sum(p.numel() for p in rnn.parameters())
assert num_params == 18, f'expected 18 parameters but found {num_params}'

rnn.load_state_dict({k: v * 0 - .1 for k, v in rnn.state_dict().items()})
data = th.FloatTensor([[[.1, .15], [.2, .25], [.3, .35], [.4, .45]], [[-.1, -1.5], [
expected_all_h = th.FloatTensor([[[[-0.1250, -0.1250, -0.1250],
    [-0.1075, -0.1075, -0.1075],
    [-0.1328, -0.1328, -0.1328],
    [-0.1452, -0.1452, -0.1452]],

    [[ 0.0600, 0.0600, 0.0600],
    [ 0.1520, 0.1520, 0.1520],
    [ 0.2344, 0.2344, 0.2344],
    [-0.0853, -0.0853, -0.0853]]])
expected_last_h = th.FloatTensor([[-0.1452, -0.1452, -0.1452],
    [-0.0853, -0.0853, -0.0853]])
all_h, last_h = rnn(data)
assert all_h.shape == expected_all_h.shape
assert last_h.shape == expected_last_h.shape
print(f'Max error all_h: {th.max(th.abs(expected_all_h - all_h)).item()}')
print(f'Max error last_h: {th.max(th.abs(expected_last_h - last_h)).item()}')

batch_size: 1
seq_len: 1
hidden_size: 1
1 torch.Size([1, 1])
torch.Size([1, 1, 1])
Expected: 0.1973753273487091, got: 0.1973753273487091, max error: 0.0
batch_size: 2
seq_len: 4
hidden_size: 3
4 torch.Size([2, 3])
torch.Size([2, 4, 3])
Max error all_h: 4.999339580535889e-05
Max error last_h: 2.498924732208252e-05

```

1.B Implementing a RNN regression model.

```
In [5]: class RecurrentRegressionModel(nn.Module):
def __init__(self, recurrent_net, output_dim=1):
    """
    Initialize a simple RNN regression model

    Inputs:
    - recurrent_net: an RNN or LSTM (single or multi layer)
    - output_dim: feature dimension of the output
    """
    super().__init__()
    self.recurrent_net = recurrent_net
    self.output_dim = output_dim
    #####
    # TODO: Initialize any parameters you need #
    # HINT: use recurrent_net.hidden_size to find the hidden state size #
    #####

    # final_layer

    # input: (batch_size, seq_len, hidden_size)

    self.final_w = nn.Linear(self.recurrent_net.hidden_size, output_dim)

    #####
    #                                     END OF YOUR CODE #
    #####

def forward(self, x):
    """
    Forward pass

    Inputs:
    - x: input tensor (B, seq_len, input_size)

    Returns:
    - out: predictions of shape (B, seq_len, self.output_dim).
    - all_h: tensor of size (B, seq_len, hidden_size) containing hidden states
            produced for each timestep.
    """
    #####
    # TODO: Implement the forward step. #
    #####

    all_h, last_h = self.recurrent_net(x)
    print("all_h shape: ", all_h.shape)

    out = self.final_w(all_h)

    # output size: (batch_size, seq_len, output_dim)
    print("out shape: ", out.shape)

    #####
    #                                     END OF YOUR CODE #
    #####

    return out, all_h
```

Tests

```
In [6]: rnn = RecurrentRegressionModel(RNNLayer(2, 3), 4)

num_params = sum(p.numel() for p in rnn.parameters())
assert num_params == 34, f'expected 34 parameters but found {num_params}'

rnn.load_state_dict({k: v * 0 - .1 for k, v in rnn.state_dict().items()})
data = th.FloatTensor([[[.1, .15], [.2, .25], [.3, .35], [.4, .45]], [[-.1, -1.5], [
expected_preds = th.FloatTensor([[[[-0.0627, -0.0627, -0.0627, -0.0627],
    [-0.0678, -0.0678, -0.0678, -0.0678],
    [-0.0604, -0.0604, -0.0604, -0.0604],
    [-0.0567, -0.0567, -0.0567, -0.0567]],

    [[-0.1180, -0.1180, -0.1180, -0.1180],
    [-0.1453, -0.1453, -0.1453, -0.1453],
    [-0.1692, -0.1692, -0.1692, -0.1692],
    [-0.0748, -0.0748, -0.0748, -0.0748]]])
expected_all_h = th.FloatTensor([[[[-0.1244, -0.1244, -0.1244],
    [-0.1073, -0.1073, -0.1073],
    [-0.1320, -0.1320, -0.1320],
    [-0.1444, -0.1444, -0.1444]],

    [[ 0.0599, 0.0599, 0.0599],
    [ 0.1509, 0.1509, 0.1509],
    [ 0.2305, 0.2305, 0.2305],
    [-0.0840, -0.0840, -0.0840]]])
preds, all_h = rnn(data)
assert all_h.shape == expected_all_h.shape
assert preds.shape == expected_preds.shape
print(f'Max error all_h: {th.max(th.abs(expected_all_h - all_h)).item()}')
print(f'Max error last_h: {th.max(th.abs(expected_preds - preds)).item()}')

batch_size: 2
seq_len: 4
hidden_size: 3
4 torch.Size([2, 3])
torch.Size([2, 4, 3])
all_h shape: torch.Size([2, 4, 3])
out shape: torch.Size([2, 4, 4])
Max error all_h: 4.699826240539551e-05
Max error last_h: 4.312396049499512e-05
```

Problem 1.C: Dataset and loss function


1.C.i: Understanding the dataset (no implementation needed)

Inspect the code and plots below to visualize the dataset

```
In [7]: def generate_batch(seq_len=10, batch_size=1):
        data = th.randn(size=(batch_size, seq_len, 1))
        sums = th.cumsum(data, dim=1)
        div = (th.arange(seq_len) + 1).unsqueeze(0).unsqueeze(2)
        target = sums / div
        return data, target
```

```
In [8]: x, y = generate_batch(seq_len=10, batch_size=4)
        for i in range(4):
            fig, ax1 = plt.subplots(1)
            ax1.plot(x[i, :, 0])
            ax1.plot(y[i, :, 0])
            ax1.legend(['x', 'y'])
            plt.title('Targets at all timesteps')
            plt.show()

        for i in range(4):
            fig, ax1 = plt.subplots(1)
            ax1.plot(x[i, :, 0])
            ax1.plot(np.arange(10), [y[i, -1].item()] * 10)
            ax1.legend(['x', 'y'])
            plt.title('Predict only at the last timestep')
            plt.show()
```



1.C.ii Implement the loss function

```
In [9]: def loss_fn(pred, y, last_timestep_only=False):
        """
        Inputs:
        - pred: model predictions of size (batch, seq_len, 1)
        - y: targets of size (batch, seq_len, 1)
        - last_timestep_only: boolean indicating whether to compute loss for all
            timesteps or only the last
        Returns:
        - loss: scalar MSE loss between pred and true labels
        """
        #####
        # TODO: implement the loss (HINT: look for pytorch's MSELoss function) #
        #####
        if last_timestep_only:
            pred = pred[:, -1]
            y = y[:, -1]
        loss_fn = nn.MSELoss()
        loss = loss_fn(pred, y)
        #####
        #                                     END OF YOUR CODE                                     #
        #####
        return loss
```

Tests

You should see errors < 1e-4

```
In [10]: pred = th.FloatTensor([[.1, .2, .3], [.4, .5, .6]])
y = th.FloatTensor([[-1.1, -1.2, -1.3], [-1.4, -1.5, -1.6]])
loss_all = loss_fn(pred, y, last_timestep_only=False)
loss_last = loss_fn(pred, y, last_timestep_only=True)
assert loss_all.shape == loss_last.shape == th.Size([])
print(f'Max error loss_all: {th.abs(loss_all - th.tensor(3.0067)).item()}')
print(f'Max error loss_last: {th.abs(loss_last - th.tensor(3.7)).item()}')
```

```
Max error loss_all: 3.314018249511719e-05
Max error loss_last: 2.384185791015625e-07
```

1.D: Analyzing RNN Gradients

You do not need to understand the details of the GradientVisualizer class in order to complete this problem.


```

In [15]: def biggest_eig_magnitude(matrix):
    """
    Inputs: a square matrix
    Returns: the scalar magnitude of the largest eigenvalue
    """
    h, w = matrix.shape
    assert h == w, f'Matrix has shape {matrix.shape}, but eigenvalues can only be computed for square matrices'
    eigs = th.linalg.eigvals(matrix)
    eig_magnitude = eigs.abs()
    eigs_sorted = sorted([i.item() for i in eig_magnitude], reverse=True)
    first_eig_magnitude = eigs_sorted[0]
    return first_eig_magnitude

class GradientVisualizer:

    def __init__(self, rnn, last_timestep_only):
        """
        Inputs:
        - rnn: rnn module
        - last_timestep_only: boolean indicating whether to compute loss for all
            timesteps or only the last timestep

        Returns:
        - loss: scalar MSE loss between pred and true labels
        """

        self.rnn = rnn
        self.last_timestep_only = last_timestep_only
        self.model = RecurrentRegressionModel(rnn)
        self.original_weights = copy.deepcopy(rnn.state_dict())

        # Generate a single batch to be used repeatedly
        self.x, self.y = generate_batch(seq_len=10)
        print(f'Data point: x={np.round(self.x[0, :, 0].detach().cpu().numpy(), 2)}, y={self.y[0]}')

    def plot_visuals(self):
        """ Generate plots which will be updated in realtime. """
        fig, (ax1, ax2) = plt.subplots(1, 2)
        ax1.set_title('RNN Outputs')
        ax1.set_xlabel('Unroll Timestep')
        ax1.set_ylabel('Hidden State Norm')
        ax1.set_ylim(-1, 5)
        plt_1 = ax1.plot(np.arange(1, 11), np.zeros(10) + 1) # placeholder vals
        plt_1 = plt_1[0]

        ax2.set_title('Gradients')
        ax2.set_xlabel('Unroll Timestep')
        ax2.set_ylabel('RNN dLoss/d a_t Gradient Magitude')
        ax2.set_ylim(10**-6, 1e5)
        ax2.set_yscale('log')
        # X-axis labels are reversed since the gradient flow is from later layers to earlier layers
        ax2.set_xticks(np.arange(10), np.arange(10, 0, -1))
        plt_2 = ax2.plot(np.arange(10), np.zeros(10) + 1) # placeholder vals
        plt_2 = plt_2[0]
        self.fig = fig
        self.plots = [plt_1, plt_2]
        return plt_1, plt_2, fig

    # Main update function for interactive plot
    def update_plots(self, weight_val=0, bias_val=0):
        # Scale the original RNN weights by a constant

```

```

w_dict = copy.deepcopy(self.original_weights)
#####
# TODO: Scale all W matrixes by weight_val, and all bias matrices by bias_val#
# If you're using PyTorch nn.Linear layers, you don't need to modify the code#
# provided, but if you're using custom layers, modify this block.          #
#####
for k in w_dict.keys():
    if 'weight' in k:
        w_dict[k][:] *= weight_val
    elif 'bias' in k:
        w_dict[k][:] *= bias_val
#####
#                                     END OF YOUR CODE                        #
#####
self.rnn.load_state_dict(w_dict)

# Don't compute for LSTMs, which don't have behavior dependent on a single eigen
if isinstance(self.rnn, RNNLayer):
    #####
    # TODO: Set W = the weight which most affects exploding/vanishing gradients #
    # Hint: Call module.weight or module.bias on the module you want to use    #
    # If you used a single Linear layer, slice a square matrix from it.        #
    #####

    # rnn.mixed_w = nn.Linear(input_size + hidden_size, hidden_size)
    # but in weight, shape of store is weight.t
    # so shape is (hidden_size, mixed_size)

    hidden_size, mixed_size = self.rnn.mixed_w.weight.shape
    # we want to get the W_h part (elim the W_x part)
    W = self.rnn.mixed_w.weight[:, -hidden_size:]
    #####
    #                                     END OF YOUR CODE                        #
    #####

    biggest_eig = biggest_eig_magnitude(W)
    print(f' Biggest eigenvalue magnitude: {biggest_eig:.3}')
```

Run model

```

pred, h = self.model(self.x)
loss = loss_fn(pred, self.y, self.last_timestep_only)
n_steps = len(h[0])

plt_1, plt_2 = self.plots

# Plot the hidden state magnitude
max_h = th.linalg.norm(h[0], dim=-1).detach().cpu().numpy()
print('Max H', ' '.join([f'{num:.3}' for num in max_h]))
plt_1.set_data(np.arange(1, n_steps + 1), np.array(max_h))
# Compute the gradient for the loss wrt the stored hidden states
# Gradients are plotted backward since we go from later layers to earlier
grads = [th.linalg.norm(num).item() for num in th.autograd.grad(loss, self.rnn.)]
print('gradients d Loss/d h_t', ' '.join([f'{num:.3}' for num in grads]))
# Add 1e-6 since it throws an error for gradients near 0
plt_2.set_data(np.arange(n_steps), np.array(grads) + 1e-6)
self.fig.canvas.draw_idle()

def create_visualization(self):
    # Include sliders for relevant quantities
    self.plot_visuals()
    ip = interactive(self.update_plots,
                     weight_val=widgets.FloatSlider(value=0, min=-5, max=5, step=.05)

```

```
        bias_val=widgets.FloatSlider(value=0, min=-5, max=5, step=.05,
    )
    return ip
```

Adjust the sliders rescale the weight and bias parameters in the RNN. Observe the effect on exploding and vanishing gradients.

Parameters to try varying:


- nonlinearity
- last_target_only

- (1) 使用no nonlinearity时, 当weight_scale过大或过小(负值)时, 梯度都会爆炸
- (2) 使用relu时, 当weight_scale过下(负值)时, 梯度会爆炸
- (3) 使用tanh时, weight_scale过大或过小(负值)时, 梯度都不会爆炸

```
In [21]: hidden_size = 16
nonlinearity = lambda x: x # options include lambda x: x (no nonlinearity), nn.fun
last_target_only = True
rnn = RNNLayer(1, hidden_size, nonlinearity=nonlinearity)
gv = GradientVisualizer(rnn, last_target_only)
gv.create_visualization()

# If for some reason the slider doesn't work for you, try calling gv.update_plots
# with various values for weight and bias
```

Data point: x=[1.51 -0.32 1.36 0.72 -0.19 0.13 -0.19 -0.43 -0.65 -0.95], y=[1.51 0.59 0.85 0.82 0.61 0.53 0.43 0.32 0.21 0.1]




```
Out[21]: interactive(children=(FloatSlider(value=0.0, description='weight_scale', layout=L
ayout(width='100%'), max=5.0, ...
```

```
In [23]: hidden_size = 16
nonlinearity = nn.functional.relu
last_target_only = True
rnn = RNNLayer(1, hidden_size, nonlinearity=nonlinearity)
gv = GradientVisualizer(rnn, last_target_only)
gv.create_visualization()

# If for some reason the slider doesn't work for you, try calling gv.update_plots
# with various values for weight and bias
```

Data point: x=[1.92 -0.3 -0.28 -0.15 -0.75 1.45 0.52 1.14 2.27 -0.41], y=[1.92 0.81 0.45 0.3 0.09 0.31 0.34 0.44 0.65 0.54]




```
Out[23]: interactive(children=(FloatSlider(value=0.0, description='weight_scale', layout=L
ayout(width='100%'), max=5.0, ...
```

```
In [27]: hidden_size = 16
nonlinearity = th.tanh
last_target_only = True
rnn = RNNLayer(1, hidden_size, nonlinearity=nonlinearity)
gv = GradientVisualizer(rnn, last_target_only)
gv.create_visualization()

# If for some reason the slider doesn't work for you, try calling gv.update_plots
# with various values for weight and bias
```

Data point: $x = [-1.68 \ 1.45 \ 2. \ -0.13 \ 0.19 \ -0.39 \ -0.31 \ -0.18 \ -0.24 \ 0.08]$, $y = [-1.68 \ -0.11 \ 0.59 \ 0.41 \ 0.37 \ 0.24 \ 0.16 \ 0.12 \ 0.08 \ 0.08]$



```
Out[27]: interactive(children=(FloatSlider(value=0.0, description='weight_scale', layout=L
ayout(width='100%'), max=5.0, ...
```

Problem 1.H: Implementing a single-layer LSTM

Hint: consider creating parameters using Pytorch's [nn.Linear](https://pytorch.org/docs/stable/generated/torch.nn.Linear.html#torch.nn.Linear) (<https://pytorch.org/docs/stable/generated/torch.nn.Linear.html#torch.nn.Linear>). You can implement this with either one Linear layer or two for each equation. If you use two, remember that you only need to include a bias term for one of the linear layers.


```

In [35]: class LSTMLayer(nn.Module):
def __init__(self, input_size, hidden_size):
    """
    Initialize a single LSTM layer.

    Inputs:
    - input_size: Data input feature dimension
    - hidden_size: RNN hidden state size (also the output feature dimension)
    """
    super().__init__()
    self.input_size = input_size
    self.hidden_size = hidden_size
    #####
    # TODO: Initialize any parameters your class needs.
    #####

    self.w = nn.Linear(input_size + hidden_size, hidden_size * 4)

    #####
    #                                     END OF YOUR CODE
    #####

def forward(self, x):
    """
    LSTM forward pass

    Inputs:
    - x: input tensor (B, seq_len, input_size)

    Returns:
    - all_h: tensor of size (B, seq_len, hidden_size) containing hidden states
      produced for each timestep
    - (h_last, c_last): hidden and cell states from the last timestep, each of
      size (B, hidden_size)
    """
    h_list = []
    #####
    # TODO: Implement the LSTM forward step
    # 1. Initialize the hidden and cell states with zeros
    # 2. Roll out the LSTM over the sequence, populating h_list along the way
    # 3. Return the appropriate outputs
    #####

    # f(t) = Sigmoid(linear1(input_size + hidden_size, hidden_size)(concat(X_t, h_t))
    # i(t) = Sigmoid(linear2(input_size + hidden_size, hidden_size)(concat(X_t, h_t))
    # o(t) = Sigmoid(linear3(input_size + hidden_size, hidden_size)(concat(X_t, h_t))
    # C(t)' = tanh(linear4((input_size + hidden_size, hidden_size)(concat(X_t, h_t))
    # C(t) = f(t) * C(t-1) + i(t) * C(t)'
    # h(t) = tanh(C(t)) * o(t)

    batch_size, seq_len = x.shape[:2]
    hs = self.hidden_size
    h_i = th.zeros((batch_size, hs)).float()
    c_i = th.zeros((batch_size, hs)).float()

    for i in range(seq_len):
        X_i = x[:, i]
        inputs = th.cat([X_i, h_i], dim = 1)
        outputs = self.w(inputs)
        #print(outputs.shape)

```

```

f_t = nn.Sigmoid()(outputs[:, :hs])
i_t = nn.Sigmoid()(outputs[:, hs:2*hs])
o_t = nn.Sigmoid()(outputs[:, 2*hs:3*hs])
c_i_hat = th.tanh(outputs[:, 3*hs:4*hs])

c_i = f_t * c_i + i_t * c_i_hat
h_i = th.tanh(c_i) * o_t

h_list.append(h_i)

h_last = h_i
c_last = c_i

#####
#                                     END OF YOUR CODE                                #
#####

# h_list should now contain all hidden states, each of size (B, hidden_size)
# We will store the hidden states so we can analyze their gradients later
self.store_h_for_grad(h_list)
all_h = th.stack(h_list, dim=1)

return all_h, (h_last, c_last)

def store_h_for_grad(self, h_list):
    """
    Store input list and allow gradient computation for all list elements
    """
    for h in h_list:
        h.retain_grad()
    self.h_list = h_list

```

Test Cases

A correct implementation should have errors < 1e-4.

```
In [36]: lstm = LSTMLayer(2, 3)
lstm.load_state_dict({k: v * 0 - .1 for k, v in lstm.state_dict().items()})
data = th.FloatTensor([[[.1, .15], [.2, .25], [.3, .35], [.4, .45]], [[-.1, -1.5], [
expected_all_h = th.FloatTensor([[-0.0273, -0.0273, -0.0273],
                                [-0.0420, -0.0420, -0.0420],
                                [-0.0514, -0.0514, -0.0514],
                                [-0.0583, -0.0583, -0.0583]],

                                [[ 0.0159,  0.0159,  0.0159],
                                [ 0.0568,  0.0568,  0.0568],
                                [ 0.1142,  0.1142,  0.1142],
                                [ 0.0369,  0.0369,  0.0369]])
expected_last_h = th.FloatTensor([[-0.0583, -0.0583, -0.0583],
                                [ 0.0369,  0.0369,  0.0369]])
expected_last_c = th.FloatTensor([[-0.1280, -0.1280, -0.1280],
                                [ 0.0759,  0.0759,  0.0759]])
all_h, (last_h, last_c) = lstm(data)
assert all_h.shape == (2, 4, 3)
assert last_h.shape == last_c.shape == (2, 3)
print(f'Max error all_h: {th.max(th.abs(expected_all_h - all_h)).item()}')
print(f'Max error last_h: {th.max(th.abs(expected_last_h - last_h)).item()}')
print(f'Max error last_h: {th.max(th.abs(expected_last_c - last_c)).item()}')
```

Max error all_h: 4.8238784074783325e-05


Max error last_h: 4.8238784074783325e-05

Max error last_h: 8.024275302886963e-06

Problem 1.8b: Analyzing gradient flow through a single-layer LSTM

```
In [37]: hidden_size = 3
last_target_only = True
rnn = LSTMLayer(1, hidden_size)
gv = GradientVisualizer(rnn, last_target_only)
gv.create_visualization()
```

Data point: $x = [1.42 \ 1.77 \ -1.32 \ -0.86 \ -0.73 \ 0.46 \ 1.88 \ 1.6 \ -0.38 \ -0.01]$, $y = [1.42 \ 1.59 \ 0.62 \ 0.25 \ 0.06 \ 0.12 \ 0.37 \ 0.53 \ 0.43 \ 0.38]$



```
Out[37]: interactive(children=(FloatSlider(value=0.0, description='weight_scale', layout=L
ayout(width='100%'), max=5.0, ...
```

Problem 1.K: Making a multi-layer RNN and LSTM

1.K.i: Implementing multi-layer models

```

In [ ]: class RNN(nn.Module):
    def __init__(self, input_size, hidden_size, num_layers):
        """
        Initialize a multilayer RNN

        Inputs:
        - input_size: Data input feature dimension
        - hidden_size: hidden state size (also the output feature dimension)
        - num_layers: number of layers
        """
        super().__init__()
        assert num_layers >= 1
        self.input_size = input_size
        self.hidden_size = hidden_size
        self.num_layers = num_layers
        #####
        # TODO: Initialize any parameters your class needs.                #
        # Consider using nn.ModuleList or nn.ModuleDict.                  #
        #####

        #####
        #                                END OF YOUR CODE                    #
        #####

    def forward(self, x):
        """
        Multilayer RNN forward pass

        Inputs:
        - x: input tensor (B, seq_len, input_size)

        Returns:
        - last_layer_h: tensor of size (B, seq_len, hidden_size) containing the
                        outputs produced for each timestep from the last layer
        - last_step_h: all hidden states from the last step (num_layers, B, hidden_size)
        """
        #####
        # TODO: Implement the RNN forward step                                #
        #####

        #####
        #                                END OF YOUR CODE                    #
        #####
        return last_layer_h, last_step_h

class LSTM(nn.Module):
    def __init__(self, input_size, hidden_size, num_layers):
        """
        Initialize a multilayer LSTM

        Inputs:
        - input_size: Data input feature dimension
        - hidden_size: hidden state size (also the output feature dimension)
        - num_layers: number of layers
        """
        super().__init__()
        assert num_layers >= 1
        self.input_size = input_size
        self.hidden_size = hidden_size
        self.num_layers = num_layers

```

```
#####
# TODO: Initialize any parameters your class needs. #
# Consider using nn.ModuleList or nn.ModuleDict. #
#####

#####
#                                     END OF YOUR CODE #
#####

def forward(self, x, hc0=None):
    """
    Multilayer LSTM forward pass

    Inputs:
    - x: input tensor (B, seq_len, input_size)

    Returns:
    - last_layer_h: tensor of size (B, seq_len, hidden_size) containing the
      outputs produced for each timestep from the last layer
    - (last_step_h, last_step_c): all hidden and cell states from the last step
      size (num_layers, B, hidden_size)
    """
    #####
    # TODO: Implement the LSTM forward step #
    #####
    #                                     END OF YOUR CODE #
    #####
    return last_layer_h, (last_step_h, last_step_c)
```


Test Cases

```

In [ ]: rnn = RNN(2, 3, 1)
rnn.load_state_dict({k: v * 0 - .1 for k, v in rnn.state_dict().items()})
data = th.FloatTensor([[[.1, .15], [.2, .25], [.3, .35], [.4, .45]], [[-.1, -1.5], [
expected_all_h = th.FloatTensor([[[[-0.1244, -0.1244, -0.1244],
                                [-0.1073, -0.1073, -0.1073],
                                [-0.1320, -0.1320, -0.1320],
                                [-0.1444, -0.1444, -0.1444]],
                                [[ 0.0599, 0.0599, 0.0599],
                                [ 0.1509, 0.1509, 0.1509],
                                [ 0.2305, 0.2305, 0.2305],
                                [-0.0840, -0.0840, -0.0840]]])
expected_last_h = th.FloatTensor([[[[-0.1444, -0.1444, -0.1444],
                                [-0.0840, -0.0840, -0.0840]]])
all_h, last_h = rnn(data)
assert all_h.shape == expected_all_h.shape
assert last_h.shape == expected_last_h.shape
print(f'Max error all_h: {th.max(th.abs(expected_all_h - all_h)).item()}')
print(f'Max error last_h: {th.max(th.abs(expected_last_h - last_h)).item()}')

rnn = RNN(2, 3, 2)
rnn.load_state_dict({k: v * 0 - .1 for k, v in rnn.state_dict().items()})
data = th.FloatTensor([[[.1, .15], [.2, .25], [.3, .35], [.4, .45]], [[-.1, -1.5], [
expected_all_h = th.FloatTensor([[[[-0.0626, -0.0626, -0.0626],
                                [-0.0490, -0.0490, -0.0490],
                                [-0.0457, -0.0457, -0.0457],
                                [-0.0430, -0.0430, -0.0430]],
                                [[[-0.1174, -0.1174, -0.1174],
                                [-0.1096, -0.1096, -0.1096],
                                [-0.1354, -0.1354, -0.1354],
                                [-0.0342, -0.0342, -0.0342]]])
expected_last_h = th.FloatTensor([[[[-0.1444, -0.1444, -0.1444],
                                [-0.0840, -0.0840, -0.0840]],
                                [[[-0.0430, -0.0430, -0.0430],
                                [-0.0342, -0.0342, -0.0342]]])
all_h, last_h = rnn(data)
assert all_h.shape == (2, 4, 3)
assert last_h.shape == (2, 2, 3)
print(f'Max error all_h: {th.max(th.abs(expected_all_h - all_h)).item()}')
print(f'Max error last_h: {th.max(th.abs(expected_last_h - last_h)).item()}')

lstm = LSTM(2, 3, 1)
lstm.load_state_dict({k: v * 0 - .1 for k, v in lstm.state_dict().items()})
data = th.FloatTensor([[[.1, .15], [.2, .25], [.3, .35], [.4, .45]], [[-.1, -1.5], [
expected_all_h = th.FloatTensor([[[[-0.0273, -0.0273, -0.0273],
                                [-0.0420, -0.0420, -0.0420],
                                [-0.0514, -0.0514, -0.0514],
                                [-0.0583, -0.0583, -0.0583]],
                                [[ 0.0159, 0.0159, 0.0159],
                                [ 0.0568, 0.0568, 0.0568],
                                [ 0.1142, 0.1142, 0.1142],
                                [ 0.0369, 0.0369, 0.0369]]])
expected_last_h = th.FloatTensor([[[[-0.0583, -0.0583, -0.0583],
                                [ 0.0369, 0.0369, 0.0369]]])
expected_last_c = th.FloatTensor([[[[-0.1280, -0.1280, -0.1280],
                                [ 0.0759, 0.0759, 0.0759]]])
all_h, (last_h, last_c) = lstm(data)
assert all_h.shape == (2, 4, 3)
assert last_h.shape == last_c.shape == (1, 2, 3)

```

```

print(f'Max error all_h: {th.max(th.abs(expected_all_h - all_h)).item()}')
print(f'Max error last_h: {th.max(th.abs(expected_last_h - last_h)).item()}')
print(f'Max error last_c: {th.max(th.abs(expected_last_c - last_c)).item()}')

lstm = LSTM(2, 3, 3)
lstm.load_state_dict({k: v * 0 - .1 for k, v in lstm.state_dict().items()})
data = th.FloatTensor([[[.1, .15], [.2, .25], [.3, .35], [.4, .45]], [[-.1, -1.5], [
expected_all_h = th.FloatTensor([[[[-0.0212, -0.0212, -0.0212],
    [-0.0296, -0.0296, -0.0296],
    [-0.0329, -0.0329, -0.0329],
    [-0.0343, -0.0343, -0.0343]],
    [[-0.0211, -0.0211, -0.0211],
    [-0.0291, -0.0291, -0.0291],
    [-0.0320, -0.0320, -0.0320],
    [-0.0332, -0.0332, -0.0332]]])
expected_last_h = th.FloatTensor([[[[-0.0583, -0.0583, -0.0583],
    [ 0.0369,  0.0369,  0.0369]],
    [[-0.0320, -0.0320, -0.0320],
    [-0.0430, -0.0430, -0.0430]],
    [[-0.0343, -0.0343, -0.0343],
    [-0.0332, -0.0332, -0.0332]]])
expected_last_c = th.FloatTensor([[[[-0.1280, -0.1280, -0.1280],
    [ 0.0759,  0.0759,  0.0759]],
    [[-0.0666, -0.0666, -0.0666],
    [-0.0907, -0.0907, -0.0907]],
    [[-0.0716, -0.0716, -0.0716],
    [-0.0693, -0.0693, -0.0693]]])
all_h, (last_h, last_c) = lstm(data)
assert all_h.shape == (2, 4, 3)
assert last_h.shape == last_c.shape == (3, 2, 3)

print(f'Max error all_h: {th.max(th.abs(expected_all_h - all_h)).item()}')
print(f'Max error last_h: {th.max(th.abs(expected_last_h - last_h)).item()}')
print(f'Max error last_c: {th.max(th.abs(expected_last_c - last_c)).item()}')

```

```

Max error all_h: 4.699826240539551e-05
Max error last_h: 4.123896360397339e-05
Max error all_h: 4.3526291847229004e-05
Max error last_h: 4.123896360397339e-05
Max error all_h: 4.8238784074783325e-05
Max error last_h: 4.8238784074783325e-05
Max error last_c: 8.024275302886963e-06
Max error all_h: 4.732981324195862e-05
Max error last_h: 4.8238784074783325e-05
Max error last_c: 4.2885541915893555e-05

```

1.K.ii: Training your model

```
In [ ]: def train(model, optimizer, num_batches, last_timestep_only, seq_len=10, batch_size=
        model
        model.train()

        losses = []
        from tqdm import tqdm
        t = tqdm(range(0, num_batches))
        for i in t:
            data, labels = generate_batch(seq_len=seq_len, batch_size=batch_size)
            pred, h = model(data)
            loss = loss_fn(pred, labels, last_timestep_only)
            losses.append(loss.item())

            optimizer.zero_grad()
            loss.backward()
            optimizer.step()
            if i % 100 == 0:
                t.set_description(f"Batch: {i} Loss: {np.mean(losses[-10:])}")
        return losses
```

```

In [ ]: def train_all(hidden_size, lr, num_batches, last_timestep_only):
    input_size = 1
    rnn_1_layer = RecurrentRegressionModel(RNN(input_size, hidden_size, 1))
    lstm_1_layer = RecurrentRegressionModel(LSTM(input_size, hidden_size, 1))
    rnn_2_layer = RecurrentRegressionModel(RNN(input_size, hidden_size, 2))
    lstm_2_layer = RecurrentRegressionModel(LSTM(input_size, hidden_size, 2))
    models = [rnn_1_layer, lstm_1_layer, rnn_2_layer, lstm_2_layer]
    model_names = ['rnn_1_layer', 'lstm_1_layer', 'rnn_2_layer', 'lstm_2_layer']

    losses = []
    for model in models:
        optimizer = optim.Adam(model.parameters(), lr=lr)
        loss = train(model, optimizer, num_batches, last_timestep_only)
        losses.append(loss)

    # visualize the results
    fig, ax1 = plt.subplots(1)
    for loss in losses:
        ax1.plot(loss)
    ax1.legend(model_names)
    plt.show()

    batch_size = 4
    x, y = generate_batch(seq_len=10, batch_size=batch_size)
    preds_list = [model(x)[0] for model in models]
    for i in range(batch_size):
        fig, ax1 = plt.subplots(1)
        ax1.plot(x[i, :, 0])
        if last_timestep_only:
            ax1.plot(np.arange(10), [y[i, -1].item()] * 10, 'bo')
        else:
            ax1.plot(y[i, :, 0], 'bo')
        for pred in preds_list:
            if last_timestep_only:
                ax1.plot(np.arange(10), [pred[i, -1, 0].detach().cpu().numpy()] * 10)
            else:
                ax1.plot(pred[i, :, 0].detach().cpu().numpy())
        ax1.legend(['x', 'y'] + model_names)
        plt.show()
    return models, losses

```

```
In [ ]: hidden_size = 32
lr = 1e-4
num_batches = 5000
last_timestep_only = False

th.manual_seed(0)
predict_all_models, predict_all_losses = train_all(hidden_size, lr, num_batches, last_timestep_only = False)
predict_one_models, predict_one_losses = train_all(hidden_size, lr, num_batches, last_timestep_only = True)
```

Batch: 4900 Loss: 0.0038075688527897: 100%|████████████████████| 5000/5000 [00:19<00:00, 254.08it/s]
Batch: 4900 Loss: 0.004596875933930278: 100%|████████████████████| 5000/5000 [00:29<00:00, 171.82it/s]
Batch: 4900 Loss: 0.0009564614854753017: 100%|████████████████████| 5000/5000 [00:19<00:00, 258.37it/s]
Batch: 4900 Loss: 0.0008792090928182005: 100%|████████████████████| 5000/5000 [00:54<00:00, 91.51it/s]

Figure

